

◀

Lesson 2:
Building a Model using SageMaker

≡

Boston Housing Example - Training the Model

SEARCH

RESOURCES

CONCEPTS

✓

1. Introduction to Amazon SageMaker

✓

2. AWS Setup Instructions for Regula...

✓

3. AWS SageMaker

✓

4. SageMaker Instance Utilization Li...

✓

5. Setting up a Notebook Instance

✓

6. Cloning the Deployment Notebooks

✓

7. Is Everything Set Up?

✓

8. Boston Housing Example - Getting...

✓

9. Boston Housing Example - Traini...

10. Boston Housing Example - Testin...

11. Mini-Project: Building Your First ...

12. Mini-Project: Solution

13. Boston Housing In-Depth - Data ...

14. Boston Housing In-Depth - Creati...

15. Boston Housing In-Depth - Buildi...

16. Boston Housing In-Depth - Creati...

17. Summary

SEND FEEDBACK



XGBoost in Competition

There's a [list of winning XGBoost-based solutions](#) to a variety of competitions, at the linked XGBoost repository.

Estimators

You can read [the documentation on estimators](#) for more information about this object. Essentially, the Estimator is an object that specifies some details about how a model will be trained. It gives you the ability to create and deploy a model.

Training Jobs

A training job is used to train a specific estimator.

When you request a training job to be executed you need to provide a few items:

- 1. A location on S3 where your training (and possibly validation) data is stored,
- 2. A location on S3 where the resulting model will be stored (this data is called the model artifacts),
- 3. A location of a docker container (certainly this is the case if using a built in algorithm) to be used for training
- 4. A description of the compute instance that should be used.

Once you provide all of this information, SageMaker will executed the necessary instance (CPU or GPU), load up the necessary docker container and execute it, passing in the location of the training data. Then when the container has finished training the model, the *model artifacts* are packaged up and stored on S3.

You can see a high-level (which we've just walked through) example of training a KMeans estimator, [in this documentation](#). This high-level example defines a KMeans estimator, and uses `.fit()` to train that model. Later, we'll show you a low-level model, in which you have to specify many more details about the training job.

Supporting Materials
[↓ XGBoost paper](#)

NEXT

